

GPU Image Filtering & Thermal Visualization

CUDA + OpenCV Capstone Project

Author: Abhinitth

Course: GPU Specialization Capstone

Project Objectives

- Demonstrate CUDA kernel programming for per-pixel operations.
- Implement GPU reductions (min/max) and normalization.
- Apply a thermal colormap LUT on GPU for visualization.
- Integrate CUDA code with OpenCV for I/O.
- Deliver reproducible build & run instructions.

Processing Pipeline

- 1. Read input image (OpenCV)**
- 2. Convert BGR → Grayscale (CUDA kernel)**
- 3. Find min/max via parallel reduction (CUDA)**
- 4. Normalize pixels to [0,1] (CUDA)**
- 5. Apply thermal colormap LUT (CUDA)**
- 6. Save output (OpenCV)**

CUDA Kernels & Implementation

Grayscale kernel uses coalesced access to read BGR and compute: $Y = 0.299 R + 0.587 G + 0.114 B$ Reduction uses shared memory per-block then global combine.
Normalization is a simple per-pixel affine transform. Colormap LUT maps normalized intensity to RGB (thermal).

Memory & Performance Considerations

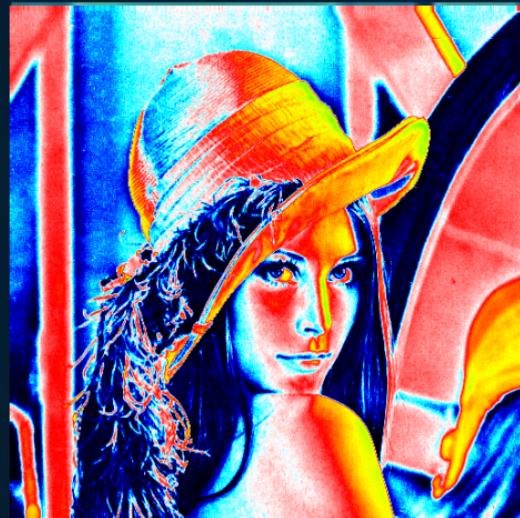
- Use global memory for image buffers (device).
- Use shared memory to speed up reductions per block.
- Minimize host-device transfers; keep work on GPU.
- Occupancy tuned with block sizes; coalescing accesses.

Parallel Reduction & Normalization

Reduction finds min and max in $O(n)$ work across threads. After reduction, normalization per pixel: $(v - \min) / (\max - \min)$. This avoids clipping and improves contrast for thermal mapping.

Thermal Colormap

**Normalized value (0..1) maps to a thermal colormap LUT on GPU. Colors progress: blue → cyan → yellow → red.
Result: visual emphasis of hot regions and gradients.**



Build & Run (Windows, VS + CUDA)

```
Open 'x64 Native Tools Command Prompt for VS 2022' cd /d  
C:\Users\abhin\OneDrive\Desktop\CUDA-Image-Processing-Capstone-Project cmake -S . -B build  
-DOpenCV_DIR="C:/Users/abhin/Downloads/opencv/build" cmake --build build --config Release cd build\Release  
GPU_Image_Filter.exe ..\..\input\input.jpg ..\..\output\output.png
```

Results & Metrics

- **Runtime:** Milliseconds for 1024x1024 images (GPU dependent)
- **Memory:** Single pass, minimal host-device transfers
- **Quality:** Thermal mapping improves visual contrast
- Presentation-ready images stored in /output folder

Future Work & Credits

- Add GPU Gaussian blur & Sobel edge detection (real-time).
- Real-time webcam thermal mode using OpenCV streaming.
- FFT-based denoising as an advanced filter option.
- Multi-GPU batching for larger datasets.

Credits: NVIDIA CUDA, OpenCV, Project Author: Abhinit